Clinical agent

the loop we keep fixed

PhysicianBench · ClinEnv

Same · tools, data, graders

LLM decides

the benchmark's own procedure

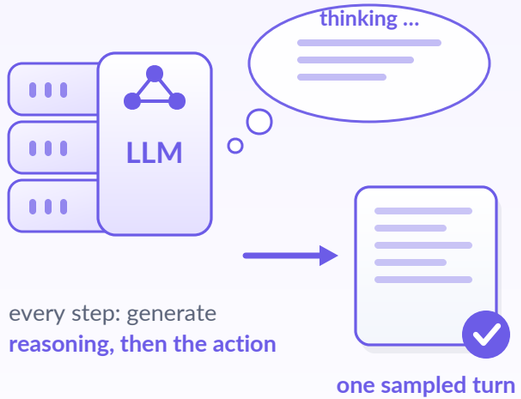

Generates every action

Thinks · thousands of tokens

System One selects

S1-Select: picks the whole turn

Decides among proposals

One pass · no generated text

System One steers

S1-Tool: picks the kind of action

Sets strategy, LLM fills in

Steers · data → action → stop