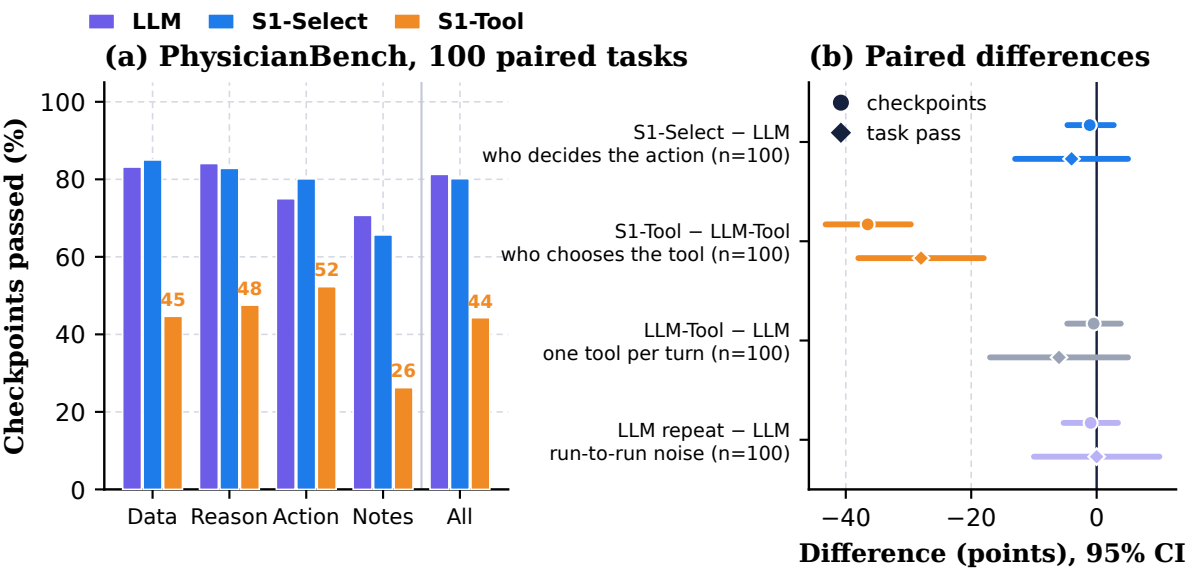
LLM
S1-Select
S1-Tool
(a) PhysicianBench, 100 paired tasks
Checkpoints passed (%)
100
80
60
40
20
0
45
48
52
26
44
Data
Reason
Action
Notes
All
(b) Paired differences
checkpoints
task pass
S1-Select − LLM
who decides the action (n=100)
S1-Tool − LLM-Tool
who chooses the tool (n=100)
LLM-Tool − LLM
one tool per turn (n=100)
LLM repeat − LLM
run-to-run noise (n=100)
−40
−20
0
Difference (points), 95% CI